

Context-quality Pareto on Long-MT-Bench+
(Hi-Seg achieves baseline-LLM-level quality at low token budget)

